

CUDA PathTracer

(Prototype)


Raimond Tunnel

Study IT in .ee




Standard Graphics Pipeline

- 1) Transform vertices (vertex shader, GPU)
- 2) Rasterize the geometry (GPU)
- 3) Color the pixels (fragment shader, GPU)



Ray Tracing


- Cast a ray (point & direction)
- Find what it hits in our scene
 - Möller-Trumbore intersection algorithm
http://en.wikipedia.org/wiki/M%C3%B6ller%E2%80%93Trumbore_intersection_algorithm
- Cast another ray from the hitpoint (in some direction)



Path Tracing

- Shoot rays from the viewer, through the pixels.
- Have them bounce around in the scene.
- Bounce with a random new direction.
- Accumulate the reflected light (BRDF).
- Upon hitting a light source, you know the color.
- Do this 800+ times per pixel, average results.

Path Tracing




Path Tracing

- Computationally more expensive than SGP.
- Random direction, in hopes of finding good enough approximation of light paths.
- Monte-Carlo like integration
- Can be easily parallelized, no dependancies between computations!

CUDA

- Language for programming Nvidia GPU-s to do general purpose computing.
(you can always program any modern GPU with shaders, but then you are restricted to SGP)
- CUDA-compatible GPU-s have a grid of blocks. Blocks have many arithmetic-logic units (ALU-s) in them and its own shared memory.

CUDA compatible GPU




Max 1024 threads
in a block.
Can queue huge
number of blocks.

My Implementation


<https://github.com/jee7/cuda-pathtracer>

- Parallelization by rays
 - Global queue of rays that need a resolution
- 1) Generate all initial rays, put them to queue
 - 2) Call a large number of worker threads to resolve the queue (dynamic parallelization)
 - 3) Save results to array and collect from CPU
 - 4) Write stuff to file, use JavaScript to see the pic.
- Because of the EENet cluster. No OpenGL or any other display.
Maybe VirtualGL?


Results (256×256, 20 iter.)




Results (256×256, 40 iter.)




Results (256×256, 80 iter.)




Results (256×256, 100 iter.)




Results (256×256, 200 iter.)




Results (256×256, 400 iter.)



Results (256×256, 600 iter.)



Results (256×256, 800 iter.)



Notice the global illumination...

Future Ideas

- Total global queue takes a lot of memory, resolve rays via batches.
- Utilize shared memory, can probably handle 1000 rays in there.
- Try out different other approaches besides the random ray bounce.

That is all... thanks!